

# AI Knowledge Assignment

---

## AI Knowledge Discovery Exercise: RAG Architecture & Unstructured Data Analysis

---

### Assignment Overview

---

**Duration:** 4-6 weeks

**Objective:** Explore how AI can leverage unstructured organizational knowledge through Retrieval-Augmented Generation (RAG) architectures and develop practical understanding of knowledge extraction from various enterprise systems.

### Learning Objectives

---

By the end of this exercise, you will:

- Understand RAG architecture fundamentals and variants
  - Build and test different RAG implementations
  - Evaluate knowledge extraction from structured and unstructured sources
  - Create a functional prototype with real-time knowledge access
  - Provide recommendations for enterprise-scale implementation
- 

## Part 1: RAG Architecture Foundation (Week 1)

---

### 1.1 Theoretical Understanding

**Goal:** Build foundational knowledge of RAG systems

#### Key Concepts to Research:

- **Retrieval-Augmented Generation (RAG):** How it combines retrieval and generation
- **Vector Embeddings:** How text is converted to numerical representations
- **Semantic Search:** Finding relevant content based on meaning, not just keywords
- **Chunk Strategies:** How to break documents into searchable pieces

- **Vector Databases:** Storage and retrieval of embeddings

### Recommended Resources:

1. Read: "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks" (original RAG paper)
2. Watch: Langchain RAG tutorials on YouTube
3. Explore: OpenAI embeddings documentation
4. Review: Pinecone/Weaviate/Chroma vector database documentation

### Deliverable:

Write a 2-page summary explaining RAG in simple terms, including diagrams showing the flow from user query to response.

## 1.2 RAG Architecture Variants

Research and document these different approaches:

### Basic RAG

- Simple retrieval → context injection → generation
- Best for: Straightforward Q&A scenarios

### Advanced RAG

- **Multi-step retrieval:** Query expansion, re-ranking
- **Fusion retrieval:** Combining multiple search strategies
- **Contextual compression:** Filtering retrieved content for relevance

### Agentic RAG

- **Tool-calling:** RAG system can decide when to search vs. generate
- **Multi-source:** Accessing different knowledge bases based on query type
- **Iterative refinement:** System can ask follow-up questions

### Deliverable:

Create a comparison matrix of RAG variants with pros/cons and use cases.

---

## Part 2: Technical Implementation (Weeks 2-3)

---

### 2.1 Environment Setup

Set up your development environment:

```
# Required tools and libraries
- Python 3.9+
- Jupyter Notebook or VS Code
- Required packages:
  * langchain
  * openai
  * chromadb (or pinecone/weaviate)
  * streamlit (for UI)
  * pandas
  * pypdf2
  * python-docx
  * requests
```

### 2.2 Build Basic RAG System

Create a simple RAG implementation:

#### Step 1: Document Processing Pipeline

```
# Your pipeline should handle:
1. PDF extraction (PyPDF2, pdfplumber)
2. Word document processing (python-docx)
3. Excel file processing (pandas)
4. Text chunking strategies:
  - Fixed size (500-1000 tokens)
  - Semantic chunking (paragraph/section breaks)
  - Sliding window with overlap
```

#### Step 2: Embedding and Storage

# Implement multiple embedding strategies:

1. OpenAI ada-002 embeddings
2. Sentence-transformers (local)
3. Compare embedding quality for different content types

### Step 3: Retrieval System

# Build retrieval with:

1. Similarity search (cosine similarity)
2. Hybrid search (keyword + semantic)
3. Metadata filtering (document type, date, department)

### Step 4: Generation Component

# Integrate with LLM:

1. OpenAI GPT models
2. Prompt engineering for context utilization
3. Citation and source tracking

## 2.3 User Interface Development

**Build a Streamlit app with:**

### File Upload Interface

- Support for PDF, DOCX, XLSX, TXT files
- Batch upload capability
- Progress indicators for processing

### RAG Configuration Panel

- Chunk size adjustment
- Embedding model selection
- Retrieval method selection
- Number of retrieved documents

### Chat Interface

- Real-time query processing

- Source document citations
- Response quality feedback

## Admin Dashboard

- Document inventory
- Search analytics
- Performance metrics

## Deliverable:

Working Streamlit application with code documentation.

---

# Part 3: Knowledge Source Integration (Week 3-4)

---

## 3.1 Structured Data Sources

Explore integration with enterprise systems:

### Salesforce Knowledge Base

- ```
# Research and prototype:
```
1. Salesforce API authentication
  2. Knowledge article extraction
  3. Metadata preservation (categories, tags, update dates)
  4. Incremental sync strategies

### SharePoint Integration

- ```
# Investigate:
```
1. SharePoint REST API / Graph API
  2. Document libraries access
  3. Permission handling
  4. Version control considerations

### Database Integration

#### # Consider:

1. FAQ databases
2. Product catalogs
3. Policy documents
4. Training materials

## 3.2 Unstructured Data Challenges

**Address real-world data issues:**

### File System Crawling

- Directory traversal strategies
- File type detection and filtering
- Duplicate document handling
- Large file processing

### Content Quality Issues

- Scanned documents (OCR requirements)
- Mixed language content
- Legacy file formats
- Corrupted or password-protected files

### Metadata Extraction

- File properties (creation date, author, department)
- Content analysis for automatic tagging
- Business context identification

## 3.3 Data Pipeline Architecture

**Design scalable ingestion:**

### Batch Processing

- Scheduled document processing
- Delta updates (only new/changed files)
- Error handling and retry logic

## Real-time Processing

- File system monitoring
- API webhooks for system updates
- Streaming data integration

### Deliverable:

Technical architecture document with data flow diagrams and integration prototypes.

---

## Part 4: Testing & Evaluation (Week 4-5)

---

### 4.1 Create Test Dataset

**Prepare diverse content for testing:**

#### Document Types

- Technical documentation
- Policy documents
- Meeting minutes
- Email threads
- Presentation slides
- Spreadsheet data

#### Test Queries

- Factual questions (direct answers in documents)
- Analytical questions (requiring synthesis)
- Comparison questions (multiple document analysis)
- Procedural questions (step-by-step processes)

### 4.2 Evaluation Metrics

**Measure system performance:**

#### Retrieval Quality

- Precision@K (relevant docs in top K results)

- Recall (relevant docs found vs. total relevant)
- MRR (Mean Reciprocal Rank)

### Generation Quality

- Answer accuracy (manual evaluation)
- Citation accuracy (sources correctly referenced)
- Response completeness
- Hallucination detection

### System Performance

- Query response time
- Document processing speed
- Storage efficiency
- Concurrent user handling

### User Experience

- Query success rate
- User satisfaction scores
- Task completion time

## 4.3 A/B Testing Framework

**Compare different approaches:**

### Chunking Strategies

- Fixed vs. semantic chunking
- Chunk size optimization
- Overlap strategies

### Retrieval Methods

- Pure semantic search
- Hybrid keyword + semantic
- Re-ranking approaches

### Generation Approaches



- Different prompt templates
- Context length optimization
- Temperature and creativity settings

### **Deliverable:**

Comprehensive evaluation report with performance benchmarks and recommendations.

---

## **Part 5: Enterprise Considerations (Week 5-6)**

---

### **5.1 Scalability Analysis**

**Assess enterprise readiness:**

#### **Technical Scalability**

- Vector database performance at scale
- Concurrent query handling
- Document processing throughput
- Storage requirements and costs

#### **Operational Scalability**

- Content governance workflows
- User access control
- System monitoring and maintenance
- Update and refresh procedures

### **5.2 Security & Compliance**

**Address enterprise concerns:**

#### **Data Security**

- Document access permissions
- Query logging and auditing
- Data encryption (at rest and in transit)
- User authentication and authorization

## Compliance Requirements

- Data retention policies
- Privacy considerations (PII handling)
- Audit trail requirements
- Regulatory compliance (if applicable)

## 5.3 Integration Strategy

### Plan for organizational adoption:

#### Change Management

- User training requirements
- Adoption strategy
- Success metrics and KPIs
- Feedback collection mechanisms

#### Technical Integration

- API design for existing systems
- Single sign-on integration
- Workflow automation possibilities
- Mobile access considerations

#### Deliverable:

Enterprise implementation roadmap with risk assessment and mitigation strategies.

---

## Part 6: Final Deliverables & Recommendations

---

### 6.1 Prototype Demo

#### Prepare comprehensive demonstration:

- Working RAG system with multiple knowledge sources
- Live query demonstration
- Performance metrics presentation
- Architecture overview

## 6.2 Technical Documentation

### Create implementation guides:

- System architecture documentation
- API documentation
- Deployment instructions
- Troubleshooting guide

## 6.3 Business Case

### Develop recommendation report:

- ROI analysis and cost considerations
- Implementation timeline
- Resource requirements
- Risk assessment
- Success metrics and measurement plan

## 6.4 Future Research Directions

### Identify next steps:

- Advanced RAG techniques to explore
- Integration opportunities
- Potential pilot programs
- Technology evolution considerations

---

## Resources & Tools

---

### Development Tools

- **Code Editor:** VS Code with Python extension
- **Notebooks:** Jupyter Lab for experimentation
- **Version Control:** Git for code management
- **Documentation:** Markdown for technical docs

### AI/ML Libraries

- **LangChain:** RAG framework and components
- **LlamaIndex:** Alternative RAG framework
- **Transformers:** Hugging Face model library
- **OpenAI:** API for embeddings and chat models

## Vector Databases

- **ChromaDB:** Local development and testing
- **Pinecone:** Cloud-based vector database
- **Weaviate:** Open-source vector database
- **FAISS:** Facebook's similarity search library

## UI Frameworks

- **Streamlit:** Rapid prototyping interface
- **Gradio:** Alternative UI framework
- **React:** For production interfaces

## Enterprise Integration

- **Salesforce APIs:** REST and GraphQL
  - **Microsoft Graph:** SharePoint and Office 365
  - **AWS/Azure:** Cloud deployment options
- 

## Success Criteria

---

### Technical Success

- ☐ Functional RAG system with multiple document types
- ☐ UI allowing file upload and real-time querying
- ☐ Integration with at least one structured data source
- ☐ Performance benchmarks and evaluation metrics
- ☐ Scalable architecture design

### Learning Success

- ☐ Clear understanding of RAG principles and variants

- ☐ Practical experience with vector embeddings and databases
- ☐ Knowledge of enterprise integration challenges
- ☐ Ability to evaluate and compare different approaches
- ☐ Business-ready recommendations and implementation plan

## Deliverable Success

- ☐ Working prototype with documentation
  - ☐ Comprehensive evaluation report
  - ☐ Enterprise implementation roadmap
  - ☐ Technical architecture documentation
  - ☐ Business case with ROI analysis
- 

## Weekly Check-ins

---

Schedule weekly progress reviews to discuss:

- Technical challenges and solutions
- Learning insights and questions
- Prototype development progress
- Enterprise considerations and requirements
- Timeline adjustments if needed

## Additional Support

---

Don't hesitate to reach out for:

- Technical guidance on implementation details
- Business context for enterprise requirements
- Access to additional resources or tools
- Clarification on objectives or deliverables

This exercise is designed to be exploratory and educational - embrace the learning process and document both successes and challenges along the way.